



(<http://vision.stanford.edu/>) (<http://stanford.edu/>)



CS231n: Convolutional Neural Networks for Visual Recognition

(<http://cs231n.stanford.edu/>)

Course Project Reports: Winter 2016

[Tweet](#)

Showing 130 out of 200 projects (70 were requested to remain private):

DeepPaint: A Tool For Image Inpainting

(http://cs231n.stanford.edu/reports/2016/pdfs/211_Report.pdf)

Suleman Kazi, Kushagr Gupta, Terry Kong



Attentional Scene Classification with Human Eye Movements

(http://cs231n.stanford.edu/reports/2016/pdfs/004_Report.pdf)

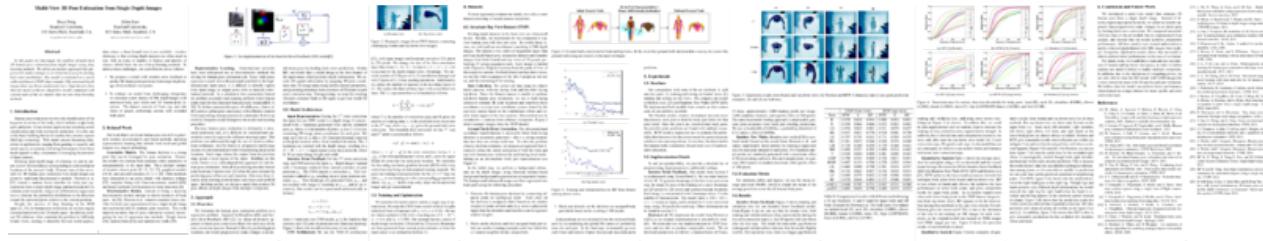
Alex Kuefler



Multi-View 3D Pose Estimation from Single Depth Images

(http://cs231n.stanford.edu/reports/2016/pdfs/412_Report.pdf)

Boya Peng, Zelun Luo



Supporting humanitarian mapping with deep learning

(http://cs231n.stanford.edu/reports/2016/pdfs/409_Report.pdf)

Lars Roemheld

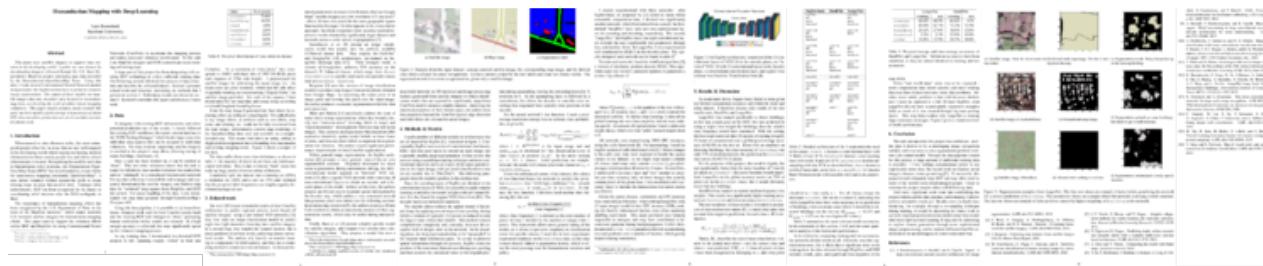
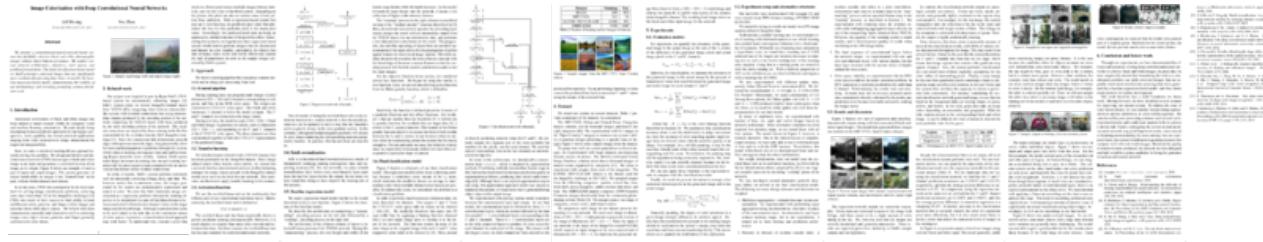


Image Colorization with Deep Convolutional Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/219_Report.pdf)

Jeff Hwang, You Zhou



The Game Imitation: A Portable Deep Learning Model for

Modern Gaming AI

(http://cs231n.stanford.edu/reports/2016/pdfs/113_Report.pdf)

Zhao Chen, Darvin Yi



Stochastic Video Prediction with Deep Conditional Generative Models

(http://cs231n.stanford.edu/reports/2016/pdfs/215_Report.pdf)

Rui Shu



Detecting Diabetic Retinopathy

(http://cs231n.stanford.edu/reports/2016/pdfs/309_Report.pdf)

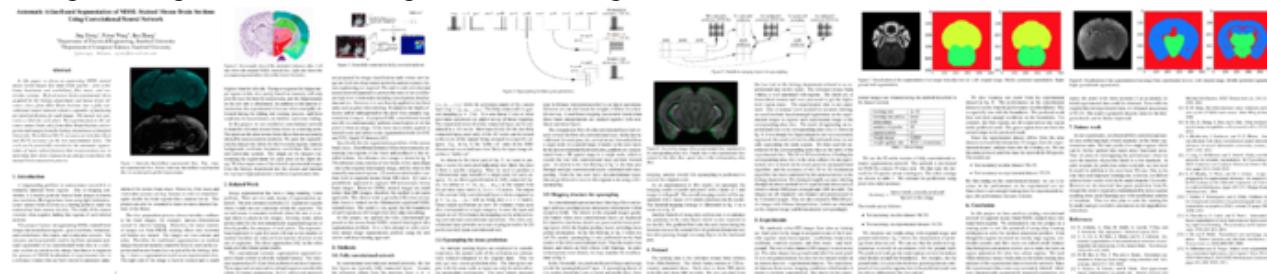
Tanner Gilligan, Marco Alban



Automatic atlas-based segmentation of NISSL stained mouse brain sections using convolutional neural network

(http://cs231n.stanford.edu/reports/2016/pdfs/318_Report.pdf)

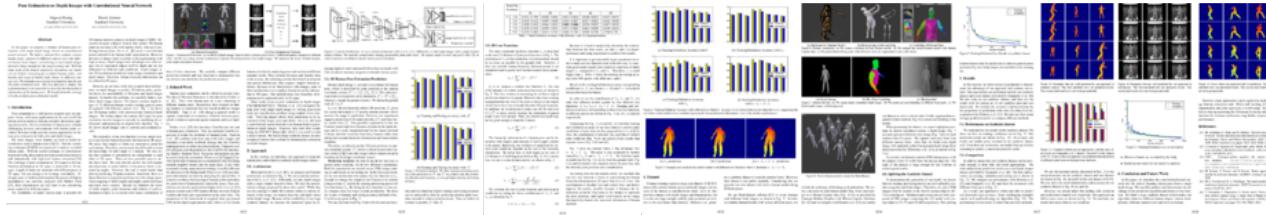
Jing Xiong, Feiran Wang, Jian Zhang



Pose Estimation on Depth Images with Convolutional Neural Network

(http://cs231n.stanford.edu/reports/2016/pdfs/426_Report.pdf)

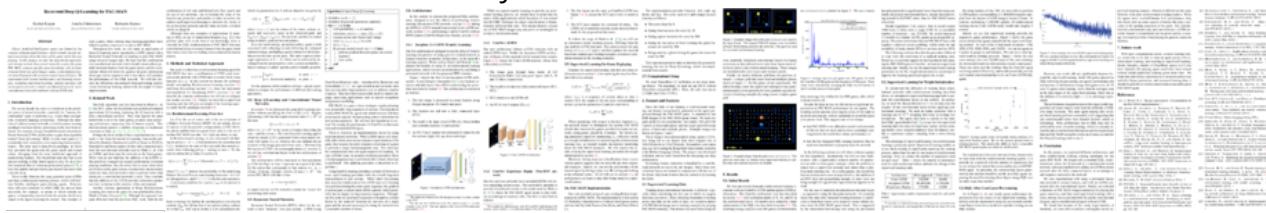
Jingwei Huang



Deep Q-Learning for PAC-MAN

(http://cs231n.stanford.edu/reports/2016/pdfs/106_Report.pdf)

Bernardo Ramos, Kushal Ranjan, Amelia Christensen



Recognizing Handwritten Characters

(http://cs231n.stanford.edu/reports/2016/pdfs/269_Report.pdf)

Lisa Yan

Automated Image Timestamp Inference using Convolutional Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/267_Report.pdf)

Prafull Sharma, Michel Shoemaker, David Pan



Show, Divide and Neural: Weighted Neural Style Transfer

(http://cs231n.stanford.edu/reports/2016/pdfs/208_Report.pdf)

Rishabh Bhargava, Ethan Chan



Neural Fill

(http://cs231n.stanford.edu/reports/2016/pdfs/209_Report.pdf)

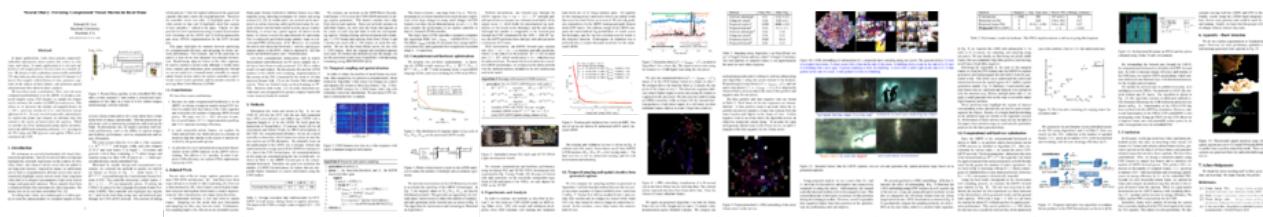
Russell Kaplan, Christopher Sauer, Alexander Lin



Neural Diary: Forming Compressed Visual Stories in Real-Time

(http://cs231n.stanford.edu/reports/2016/pdfs/213_Report.pdf)

Edward Lee



Model-Based Reinforcement Learning for Playing Atari Games

(http://cs231n.stanford.edu/reports/2016/pdfs/116_Report.pdf)

Irving Hsu, Justin Fu



Deep Reinforcement Learning for Simulated Autonomous Vehicle Control

(http://cs231n.stanford.edu/reports/2016/pdfs/112_Report.pdf)

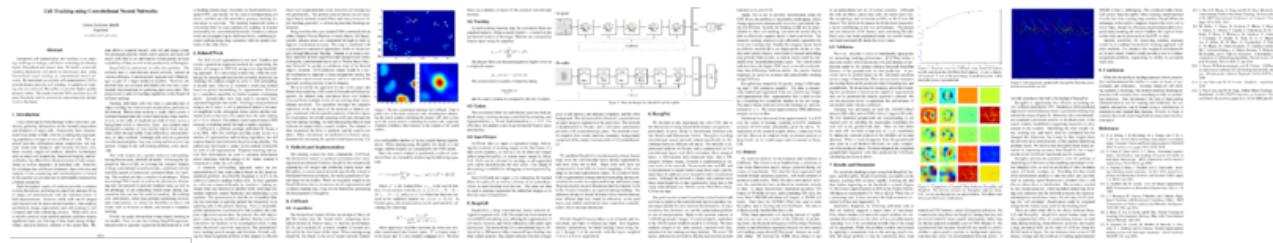
Raphael Palefsky-Smith, April Yu, Rishi Bedi



Cell Nucleus Tracking using Convolutional Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/326_Report.pdf)

Anton Jackson-Smith



Optimizing Convolutional Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/108_Report.pdf)

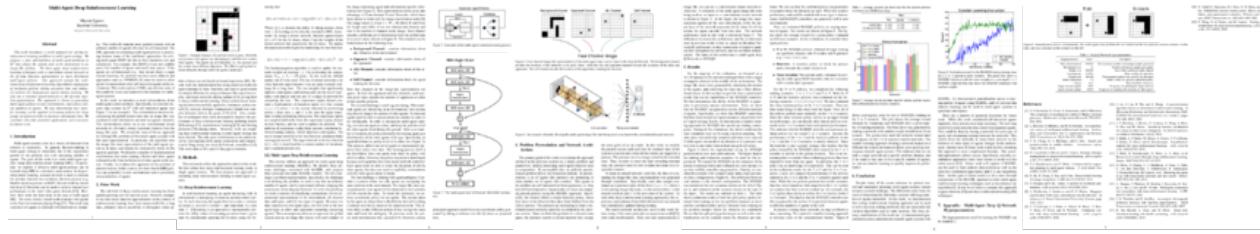
Timothy Dozat



Multi-Agent Deep Reinforcement Learning

(http://cs231n.stanford.edu/reports/2016/pdfs/122_Report.pdf)

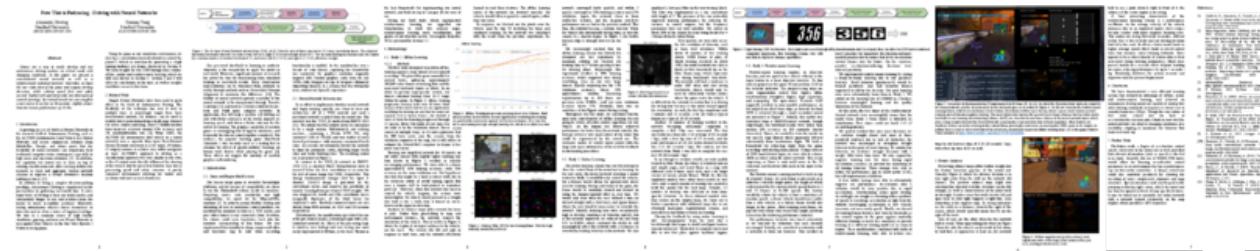
Maxim Egorov



Now This is Pod Racing

(http://cs231n.stanford.edu/reports/2016/pdfs/100_Report.pdf)

Alexander Dewing, Xiaonan Tong



Deep Learning of Spatial and Temporal Features for Automotive Prediction

(http://cs231n.stanford.edu/reports/2016/pdfs/102_Report.pdf)

Tim Wheeler, Jeremy Morton



Complexity Beyond the Trigram: Identifying Sign Languages from Video

(http://cs231n.stanford.edu/reports/2016/pdfs/207_Report.pdf)

Pamela Toman



Signature Verification with Convolutional Neural Networks (http://cs231n.stanford.edu/reports/2016/pdfs/276_Report.pdf)

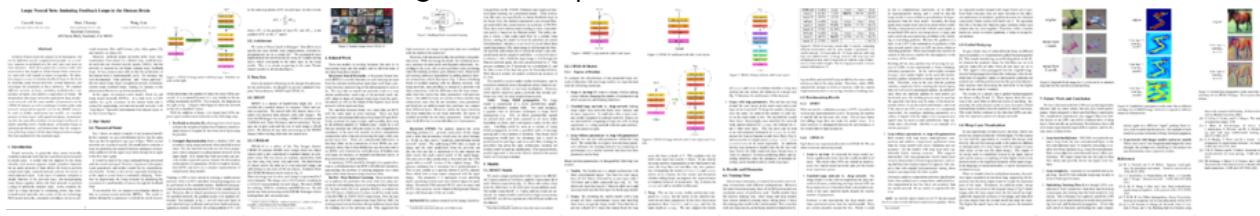
Gabriel Alvarez, Blue Sheffer, Morgan Bryant



Loopy Neural Nets: Imitating Feedback Loops in the Human Brain

(http://cs231n.stanford.edu/reports/2016/pdfs/110_Report.pdf)

Isaac Caswell, Lisa Wang, Chuanqi Shen



DeepFace: Face Generation using Deep Learning

(http://cs231n.stanford.edu/reports/2016/pdfs/006_Report.pdf)

Hardie Cate, Fahim Dalvi, Zeshan Hussain



Pruning of Winograd and FFT Convolution Algorithm

(http://cs231n.stanford.edu/reports/2016/pdfs/117_Report.pdf)

Xingyu Liu, Yatish Turakhia



Playing Super Hexagon with Convolutional Neural Networks
[\(http://cs231n.stanford.edu/reports/2016/pdfs/115_Report.pdf\)](http://cs231n.stanford.edu/reports/2016/pdfs/115_Report.pdf)
 Jason Lewis



Heirarchical Deep Convolutional Neural Networks: An Insect Classification Application
[\(http://cs231n.stanford.edu/reports/2016/pdfs/283_Report.pdf\)](http://cs231n.stanford.edu/reports/2016/pdfs/283_Report.pdf)
 Jeffrey Glick, Katarina Miller



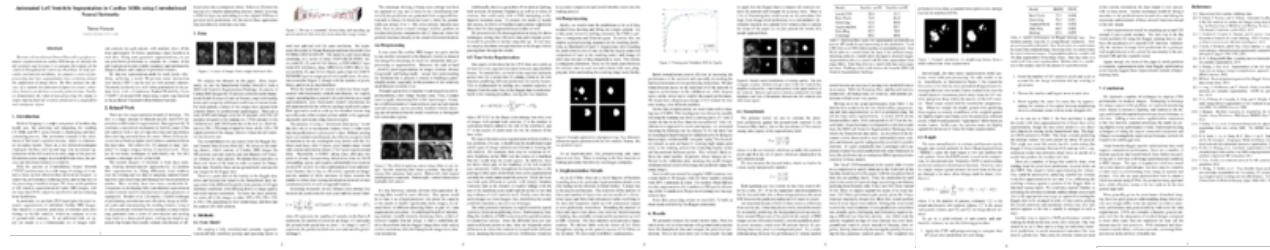
Convolutional Neural Networks for Left Ventricle Volume Estimation
[\(http://cs231n.stanford.edu/reports/2016/pdfs/305_Report.pdf\)](http://cs231n.stanford.edu/reports/2016/pdfs/305_Report.pdf)
 Cheryl Danner, Carol Hsin



Automated Left Ventricle Segmentation in Cardiac MRIs with Convolutional Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/317_Report.pdf)

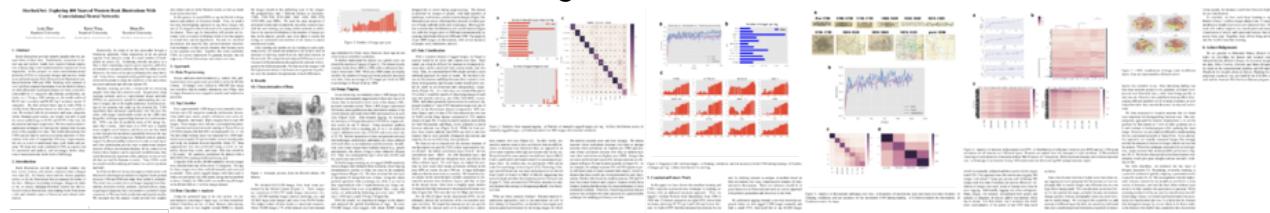
Taman Narayan



SherlockNet: Exploring 100 Years of Western Book Illustrations

(http://cs231n.stanford.edu/reports/2016/pdfs/204_Report.pdf)

Brian Do, Luda Zhao, Karen Wang



Automated Bone Age Assessment

(http://cs231n.stanford.edu/reports/2016/pdfs/310_Report.pdf)

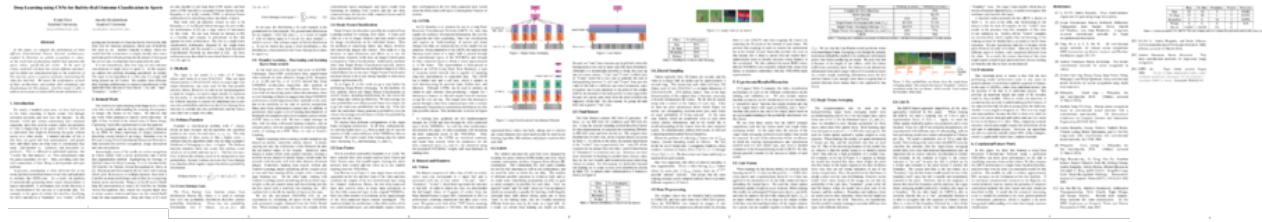
Matthew Chen



Ball-by-Ball Outcome Prediction for Cricket Matches

(http://cs231n.stanford.edu/reports/2016/pdfs/273_Report.pdf)

Anusha Balakrishnan, Kalpit Dixit



Object Detection with Reduced Image Qualities

(http://cs231n.stanford.edu/reports/2016/pdfs/287_Report.pdf)

Shiquan Wang, Hao Jiang



Pulmonary Nodule Classification with Convolutional Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/324_Report.pdf)

Sheila Ramaswamy, Karen Truong



DeepVideo: Video Summarization using Temporal Sequence Modelling

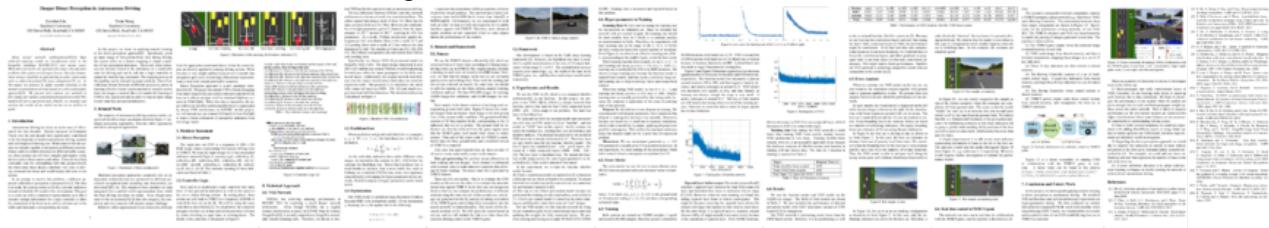
(http://cs231n.stanford.edu/reports/2016/pdfs/216_Report.pdf)

Juhi Amitkumar Naik



Deep Residual Learning for Direct Perception in Autonomous Driving

(http://cs231n.stanford.edu/reports/2016/pdfs/123_Report.pdf)
Zuozhen Liu, Yixin Wang



Distilling Knowledge to Specialist ConvNets for Object Recognition with Clustered Classes

(http://cs231n.stanford.edu/reports/2016/pdfs/120_Report.pdf)
Nathanael Romano, Robin Schucker



Measuring Cellular Semantic Distance with Siamese Neural Nets

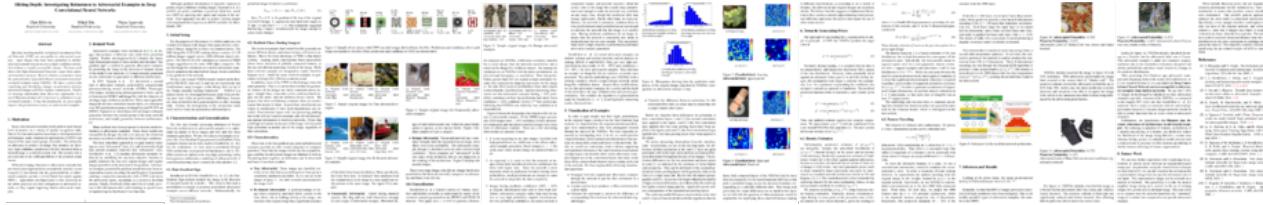
(http://cs231n.stanford.edu/reports/2016/pdfs/320_Report.pdf)
Morgan Paull



Hitting Depth: Investigating Robustness to Adversarial Examples in Deep Convolutional Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/119_Report.pdf)

Mihail Eric, Chris Billovits, Nipun Agarwala



Using Convolutional and Recurrent Neural Networks on
Magnetoencephalography Data to Assess Neuroplasticity
(http://cs231n.stanford.edu/reports/2016/pdfs/334_Report.pdf)
Irán Román



Automatic Colorization with Deep Convolutional Generative
Adversarial Networks
(http://cs231n.stanford.edu/reports/2016/pdfs/224_Report.pdf)
Stephen Koo



Disney Characters: Classifying Commercial Merchandise
based on Cartoon Characters
(http://cs231n.stanford.edu/reports/2016/pdfs/265_Report.pdf)
Paula Kusumaputri, Edwin Park

This poster presents a project titled "Classifying Objects Using Convolutional Neural Networks". It includes sections on Introduction, Related Work, Dataset, Model Architecture, Training, Evaluation, and Conclusion. The poster features several figures, including a grid of flags from various countries, a scatter plot of accuracy vs. learning rate, and a heatmap of confusion matrices.

DeepMammo

(http://cs231n.stanford.edu/reports/2016/pdfs/306_Report.pdf)

Daniel Levy, Arzav Jain

This poster discusses augmenting nearest neighbor-based algorithms with Siamese neural networks. It includes sections on Introduction, Related Work, Methodology, Experiments, and Conclusion. The poster contains numerous figures, such as line graphs showing performance metrics like accuracy and loss over time, and visualizations of feature maps and heatmaps.

Augmenting Nearest Neighbor-Based Algorithms with Siamese Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/258_Report.pdf)

Gene Lewis, Wilbur Yang

This poster details a traffic sign detection system using the You Only Look Once (YOLO) framework. It covers the dataset, model architecture, training, evaluation, and results. The poster includes many visual examples of detected traffic signs and corresponding heatmaps.

Traffic Sign Detection Using You Only Look Once Framework

(http://cs231n.stanford.edu/reports/2016/pdfs/263_Report.pdf)

Royce Cheng-Yue, Chung Yu Wang

This poster provides a comprehensive overview of traffic sign detection using the YOLO framework. It includes sections on the dataset, model architecture, training, and experimental results. The poster features many images of detected traffic signs and their bounding boxes, along with quantitative performance metrics.

ResFuse and RefiNets: Enhanced CNN architectures for image classification

(http://cs231n.stanford.edu/reports/2016/pdfs/277_Report.pdf)

Frederic Ren, Allen Nie, Nihit Desai



Predict attribute labels for restaurants using user-submitted photos

(http://cs231n.stanford.edu/reports/2016/pdfs/025_Report.pdf)

Shubham Gupta, Vinaya Polamreddi



Using Convolutional Neural Networks to demystify aesthetic works of art

(http://cs231n.stanford.edu/reports/2016/pdfs/210_Report.pdf)

Pujun Bhatnagar



Let Blind People See: Real-Time Visual Recognition with Results Converted to 3-D Audio

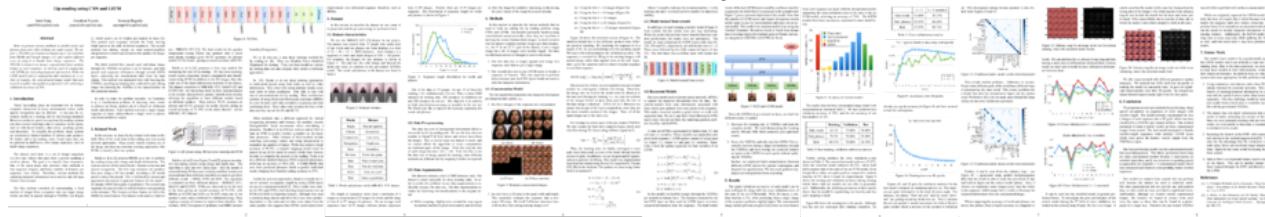
(http://cs231n.stanford.edu/reports/2016/pdfs/218_Report.pdf)

Rui (Forest) Jiang, Qian Lin, Shuhui Qu



Lip-Reading with CNNs and LSTMs

(http://cs231n.stanford.edu/reports/2016/pdfs/217_Report.pdf)
Sameep Bagadia, Amit Garg, Jonathan Noyola



Tracking a Parking Lot's Empty Spaces Without Sensors

(http://cs231n.stanford.edu/reports/2016/pdfs/280_Report.pdf)
Jordan Cazamias, Martina Marek



Detecting and Classifying Lung Nodules in Chest Radiographs

(http://cs231n.stanford.edu/reports/2016/pdfs/313_Report.pdf)
Isabel Bush



Multi-Frame Video Super Resolution using Deep Convolutional Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/212_Report.pdf)

Alex Greaves, Hanna Winter



Facial Affect Detection Using Convolutional Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/011_Report.pdf)

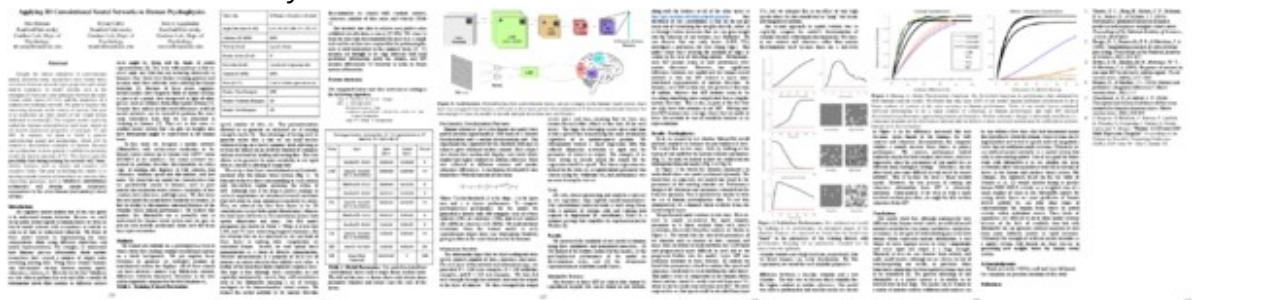
Sherrie Wang



Solving human psychophysics tasks with 3D convolutional neural networks

(http://cs231n.stanford.edu/reports/2016/pdfs/408_Report.pdf)

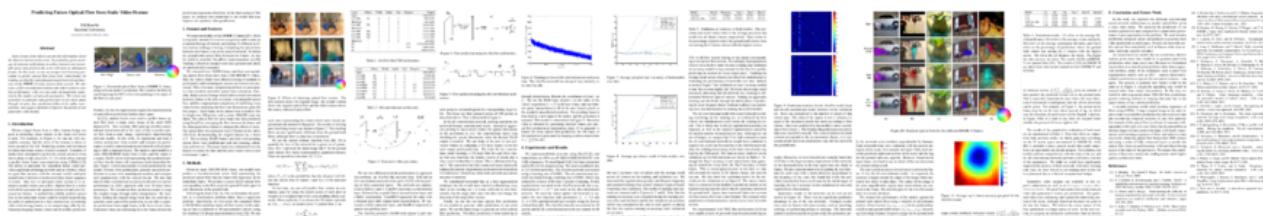
Daniel Birman, Dylan Cable



Predicting future optical flow from static video frames

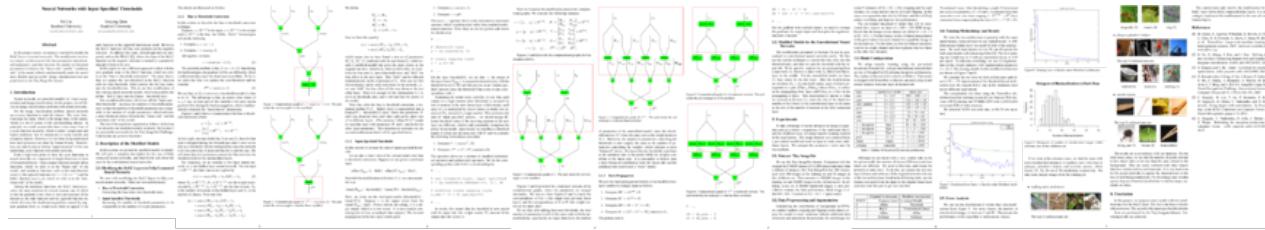
(http://cs231n.stanford.edu/reports/2016/pdfs/223_Report.pdf)

Pol Rosello



Neural Networks with Input Specified Thresholds (http://cs231n.stanford.edu/reports/2016/pdfs/118_Report.pdf)

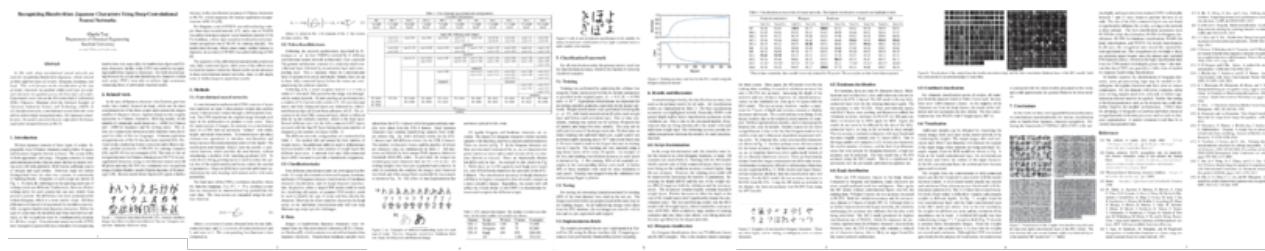
Fei Liu, Junyang Qian



Recognizing Handwritten Japanese Characters Using Deep Convolutional Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/262_Report.pdf)

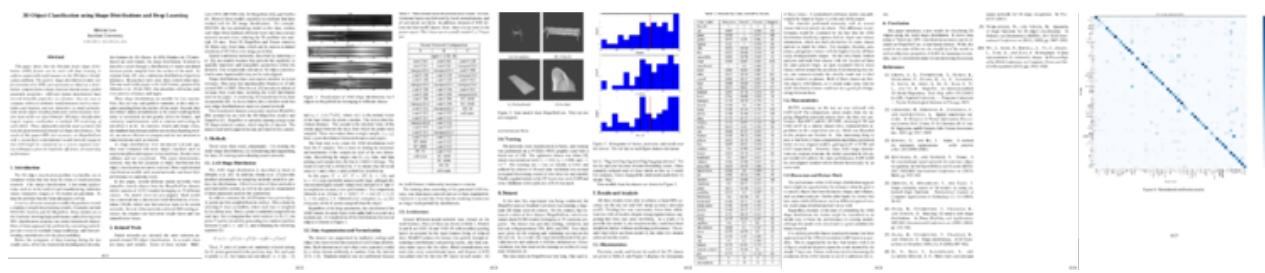
Charlie Tsai



3D Object Classification using Shape Distributions and Deep Learning

(http://cs231n.stanford.edu/reports/2016/pdfs/420_Report.pdf)

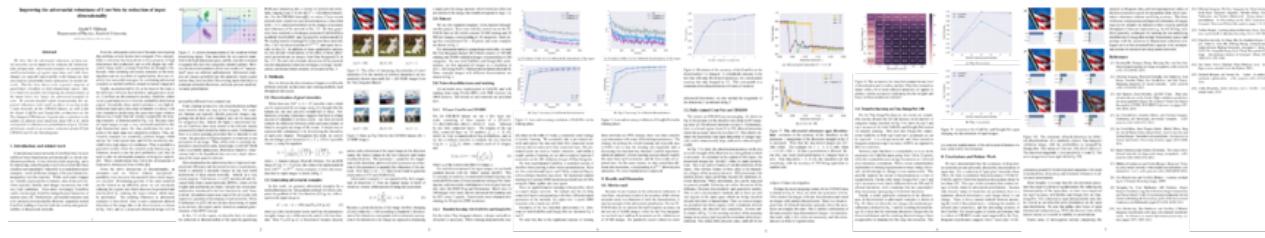
Melvin Low



Improving the adversarial robustness of Convnets

(http://cs231n.stanford.edu/reports/2016/pdfs/103_Report.pdf)

Akash Maharaj



Playing Flappybird with Deep Reinforcement Learning

(http://cs231n.stanford.edu/reports/2016/pdfs/111_Report.pdf)

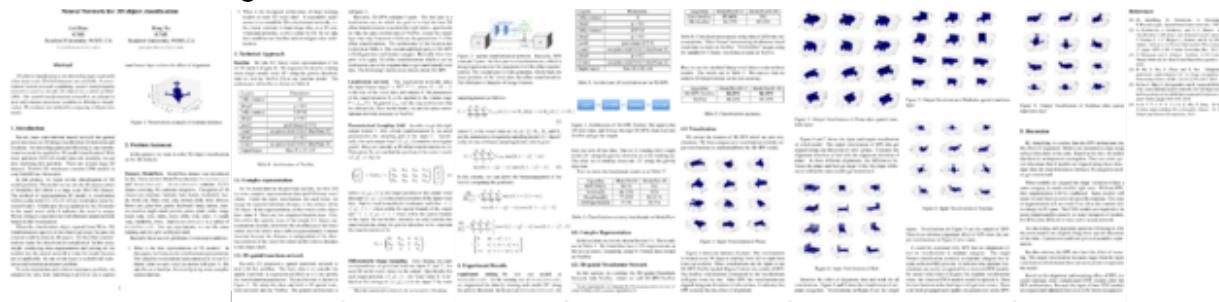
Naveen Appiah Balaji, Sagar Vare



Neural Network for 3D object classification

(http://cs231n.stanford.edu/reports/2016/pdfs/417_Report.pdf)

Lin Shao, Peng Xu



Facial Expression Recognition with Convolutional Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/023_Report.pdf)

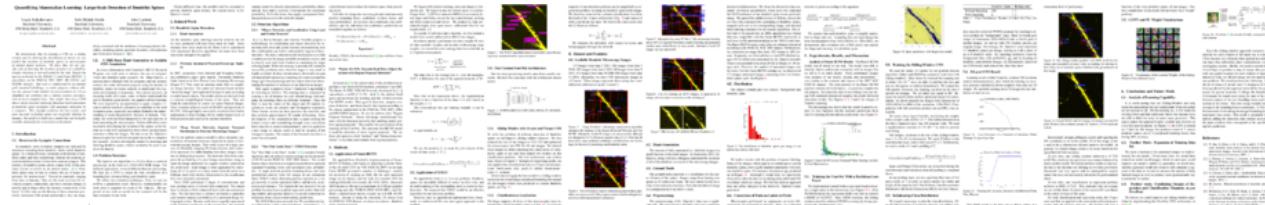
Vivek Choksi, Arushi Raghuvanshi



Quantifying Mammalian Learning: Large-Scale Detection of Dendritic Spines

(http://cs231n.stanford.edu/reports/2016/pdfs/329_Report.pdf)

Seth Hildick-Smith, Ivaylo Bahtchevanov , John Lambert



Intelligent coloring of greyscale images using CNNs

(http://cs231n.stanford.edu/reports/2016/pdfs/202_Report.pdf)

Tian Zhao, Ross Daly



Convolutional Neural Networks for Facial Age and Gender Prediction

(http://cs231n.stanford.edu/reports/2016/pdfs/003_Report.pdf)

Ari Ekmekji



ambience_is_classy or good_for_kids? : Yelp Restaurant Classification Challenge

(http://cs231n.stanford.edu/reports/2016/pdfs/015_Report.pdf)

Jee Ian Tam



Understanding Visual Art with CNNs

(http://cs231n.stanford.edu/reports/2016/pdfs/201_Report.pdf)

Michael Baumer



Action Recognition and Detection Using CNNs and LSTMs

(http://cs231n.stanford.edu/reports/2016/pdfs/221_Report.pdf)

Helen Jiang, Gary Thung



Facial keypoints detection using Neural Network

(http://cs231n.stanford.edu/reports/2016/pdfs/007_Report.pdf)

Chenyue Meng, Shutong Zhang



Near Real-Time Super-Resolution Video

(http://cs231n.stanford.edu/reports/2016/pdfs/203_Report.pdf)

David Zeng



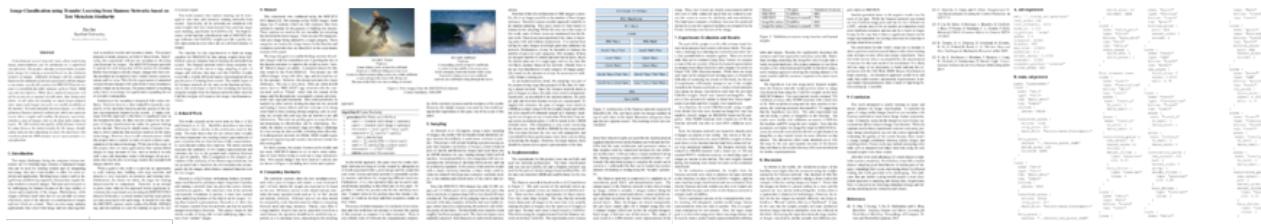
Transforming How We Diagnose Heart Disease Using Convolutional Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/327_Report.pdf)
Viswajith Venugopal, Swaroop Ramaswamy



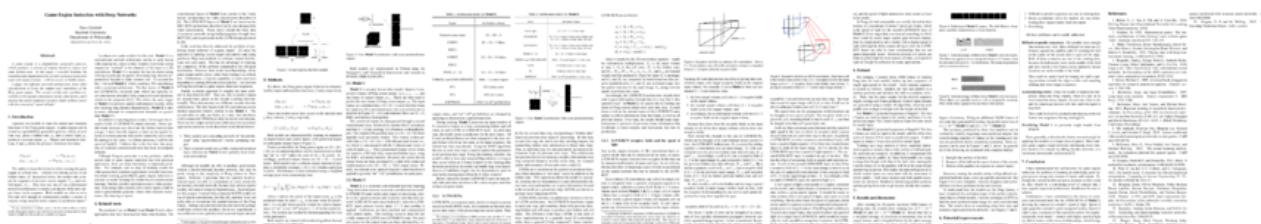
Improving image feature representations with noisy similarity of text captions

(http://cs231n.stanford.edu/reports/2016/pdfs/360_Report.pdf)
Dan Iter



A Recurrent Neural Game Engine

(http://cs231n.stanford.edu/reports/2016/pdfs/101_Report.pdf)
David Gottlieb



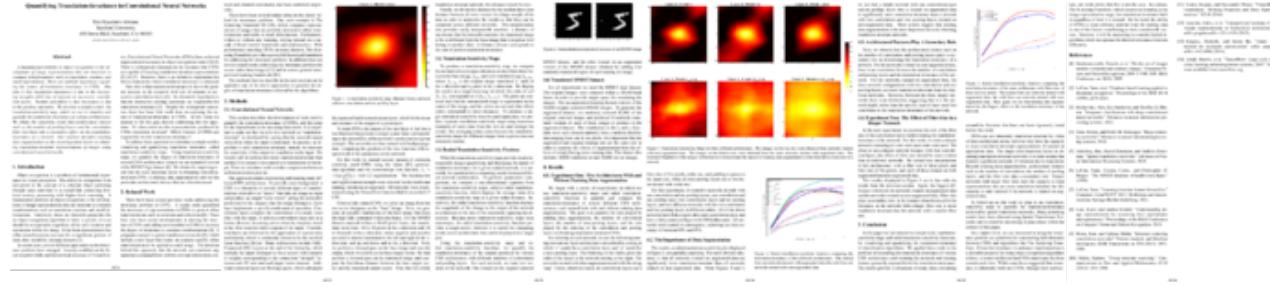
Playing Tetris with Deep Reinforcement Learning (http://cs231n.stanford.edu/reports/2016/pdfs/121_Report.pdf)

Matt Stevens, Sabeek Pradhan



Quantifying Invariance in Convolutional Neural Networks (http://cs231n.stanford.edu/reports/2016/pdfs/107_Report.pdf)

Eric Kauderer-Abrams



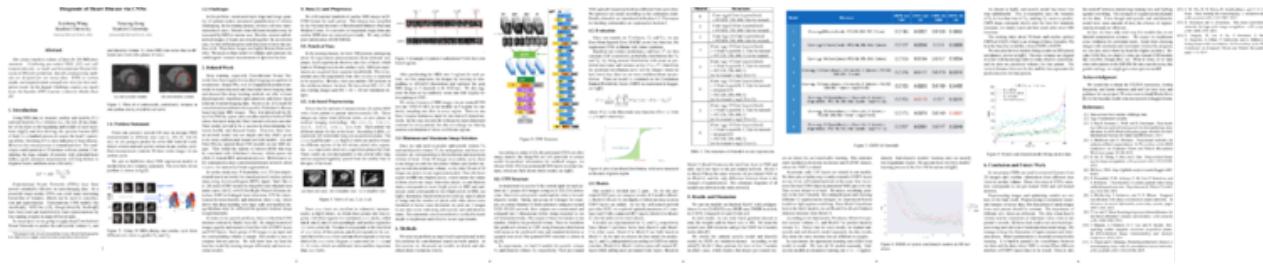
Yelp Restaurant Photo Classification (http://cs231n.stanford.edu/reports/2016/pdfs/013_Report.pdf)

Jade Huang



Diagnosis of Heart Disease via CNNs (http://cs231n.stanford.edu/reports/2016/pdfs/331_Report.pdf)

Yanyang Kong, Kaicheng Wang



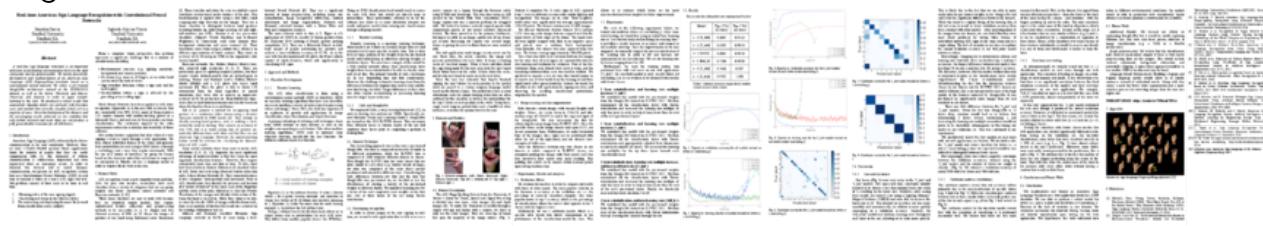
Single Image Depth Prediction Using Convolutional Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/407_Report.pdf)
Ralph Ma, Jack Zhu



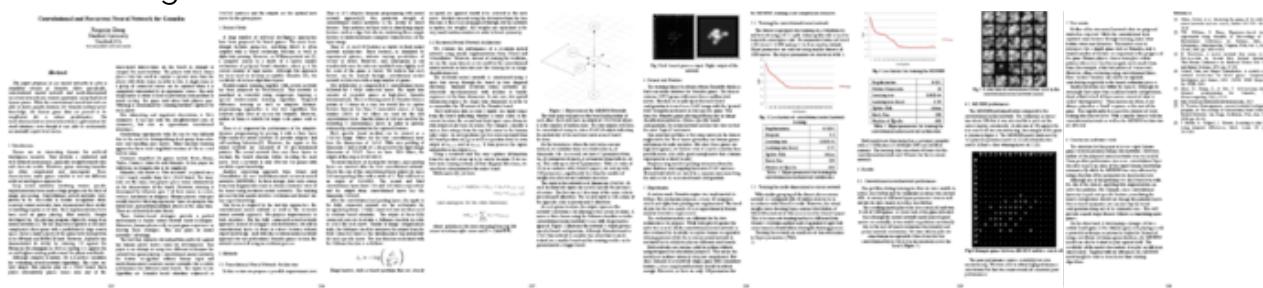
Read my Hands: Translating American Sign Language to Text

(http://cs231n.stanford.edu/reports/2016/pdfs/214_Report.pdf)
Sigberto Alarcon Viesca, Brandon Garcia



Game playing using recurrent neural networks

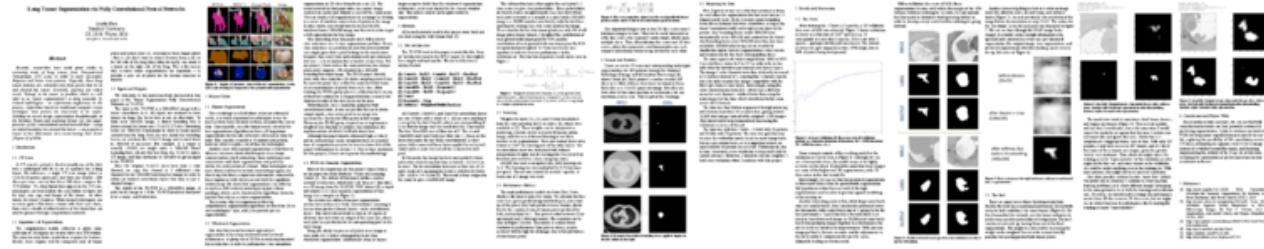
(http://cs231n.stanford.edu/reports/2016/pdfs/109_Report.pdf)
Nic Zhang



Fully Connected Convolutional Neural Networks for Lung Tumor Segmentation

(http://cs231n.stanford.edu/reports/2016/pdfs/302_Report.pdf)

Austin Ray



Residual Networks for Tiny ImageNet

(http://cs231n.stanford.edu/reports/2016/pdfs/411_Report.pdf)

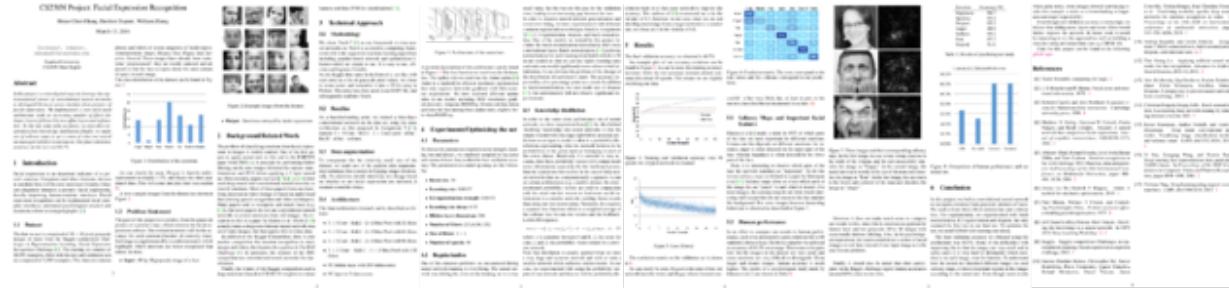
Hansohl Kim



Facial Expression Recognition

(http://cs231n.stanford.edu/reports/2016/pdfs/009_Report.pdf)

Hsiao Chen Chang, Emilien Dupont, William Zhang



Geolocalization of Street View Images: From Pixels to Places

(http://cs231n.stanford.edu/reports/2016/pdfs/422_Report.pdf)

Blake Wulfe, David Hershey



Classifying Restaurant Attributes With Deep Learning

(http://cs231n.stanford.edu/reports/2016/pdfs/001_Report.pdf)

Pedro Garzon, Diveesh Singh



Abstract concept & emotion detection of tagged images with CNNs

(http://cs231n.stanford.edu/reports/2016/pdfs/008_Report.pdf)

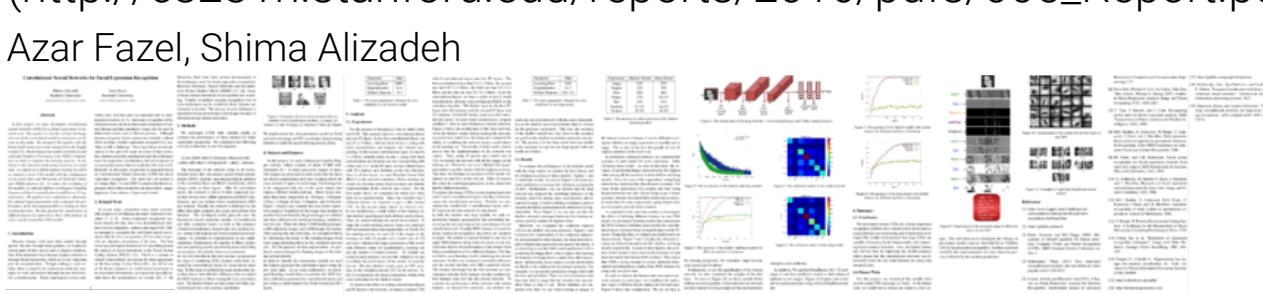
Youssef Ahres, Nikolaus Volk



Convolutional Neural Networks for Facial Expression Recognition

(http://cs231n.stanford.edu/reports/2016/pdfs/005_Report.pdf)

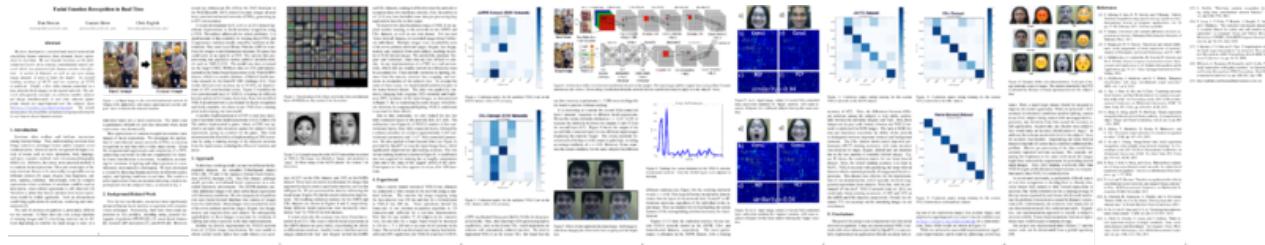
Azar Fazel, Shima Alizadeh



Facial Emotion Recognition in Real Time

(http://cs231n.stanford.edu/reports/2016/pdfs/022_Report.pdf)

Dan Duncan, Chris English, Gautam Shine

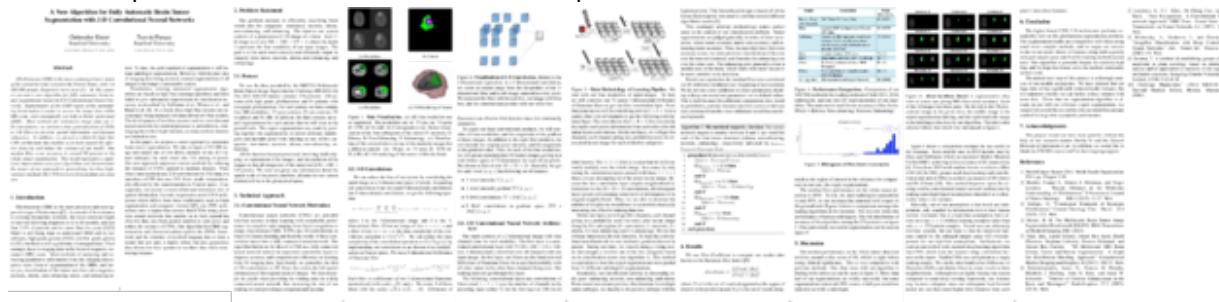


A New Algorithm for Fully Automatic Brain Tumor

Segmentation with 3-D Convolutional Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/322_Report.pdf)

Christopher Elamri, Teun de Planque



Fast Unsupervised Object Localization

(http://cs231n.stanford.edu/reports/2016/pdfs/285_Report.pdf)

Mihir Mongia, deepak menghani, Anjan Dwaraknath



Facial Keypoint Detection

(http://cs231n.stanford.edu/reports/2016/pdfs/010_Report.pdf)

Shayne Longpre, Ajay Sohmshetty

This block contains a collage of 12 pages from a report. The pages are organized into three columns and four rows. The first column contains pages 1 through 4. The second column contains pages 5 through 8. The third column contains pages 9 through 12. The pages include various sections such as 'Introduction', 'Method', 'Results', and 'Conclusion', along with tables, figures, and diagrams.

Cardiovascular Edge Detection using Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/311_Report.pdf)

Gabriel Maher

This block contains a collage of 12 pages from a report. The pages are organized into three columns and four rows. The first column contains pages 1 through 4. The second column contains pages 5 through 8. The third column contains pages 9 through 12. The pages include various sections such as 'Introduction', 'Method', 'Results', and 'Conclusion', along with tables, figures, and diagrams.

Convolution Neural Networks for Chinese Handwriting Recognition

(http://cs231n.stanford.edu/reports/2016/pdfs/428_Report.pdf)

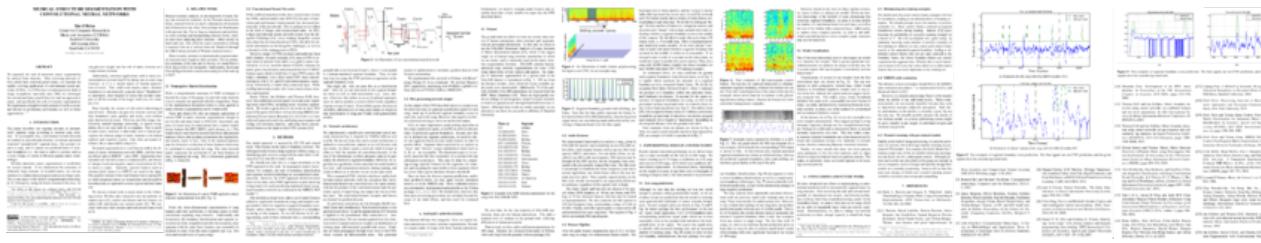
Xu Chen

This block contains a collage of 12 pages from a report. The pages are organized into three columns and four rows. The first column contains pages 1 through 4. The second column contains pages 5 through 8. The third column contains pages 9 through 12. The pages include various sections such as 'Introduction', 'Method', 'Results', and 'Conclusion', along with tables, figures, and diagrams.

Music Structure Segmentation with Convolutional Neural Networks

(http://cs231n.stanford.edu/reports/2016/pdfs/220_Report.pdf)

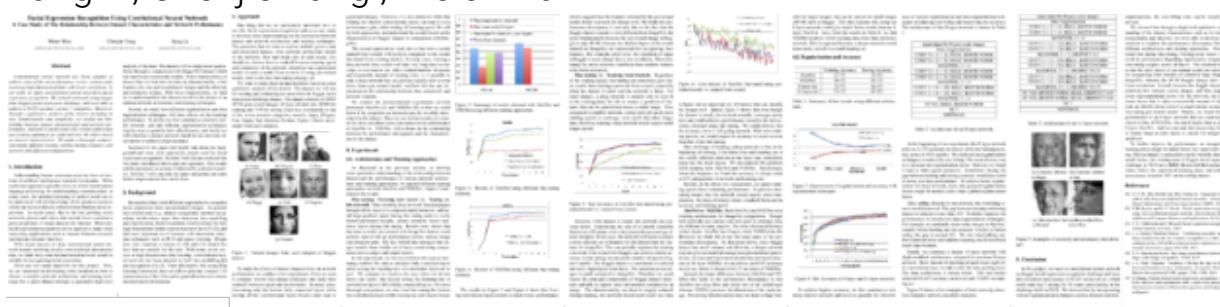
Tim O'Brien



Facial Expression Recognition with Selected Partial Facial Feature Using CNN

(http://cs231n.stanford.edu/reports/2016/pdfs/020_Report.pdf)

Yang Li, Chenjie Yang , Weier Wan



Brain tumour segmentation from MRI

(http://cs231n.stanford.edu/reports/2016/pdfs/328_Report.pdf)

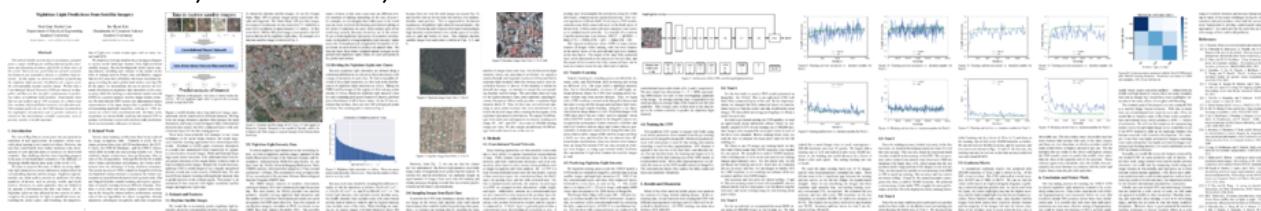
Raunaq Rewari



Nighttime Light Predictions from Satellite Imagery

(http://cs231n.stanford.edu/reports/2016/pdfs/423_Report.pdf)

Rachel Luo, Neal Jean, Jake Kim



Coloring Grayscale image

(http://cs231n.stanford.edu/reports/2016/pdfs/205_Report.pdf)

Vivek Bagaria, Kedar Tatwawadi



Yelp Restaurant Photo Classification: MIML Approach with CNNs

(http://cs231n.stanford.edu/reports/2016/pdfs/017_Report.pdf)

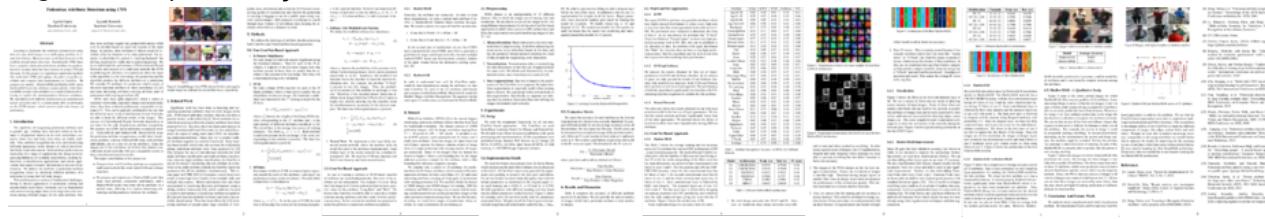
Raghav Gupta, Pulkit Agrawal



Pedestrian Attribute Detection using CNN

(http://cs231n.stanford.edu/reports/2016/pdfs/255_Report.pdf)

Agrim Gupta, Jayanth Ramesh



Multiple Object Recognition with Focusing and Blurring

(http://cs231n.stanford.edu/reports/2016/pdfs/259_Report.pdf)

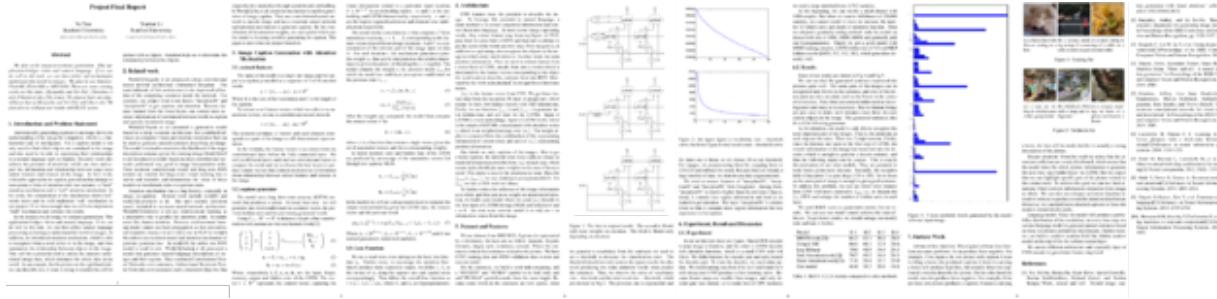
Holly Chiang, Yifan Ge, Connie Wu



Image Captioning with Attention

(http://cs231n.stanford.edu/reports/2016/pdfs/364_Report.pdf)

Ye Tian, Tianlun Li



Machine Cashier

(http://cs231n.stanford.edu/reports/2016/pdfs/260_Report.pdf)

Roy Chan



Tractable Neural Networks for Identity Recognition

(http://cs231n.stanford.edu/reports/2016/pdfs/250_Report.pdf)

David Eng, Jenny Hong



Models of Attention for Video Captioning

(http://cs231n.stanford.edu/reports/2016/pdfs/362_Report.pdf)

Dieterich Lawson, Blaine Rister

This report page contains a grid of sections and figures. It includes a title, abstract, introduction, related work, methodology, experiments, results, and conclusion. There are also several small images and plots illustrating the geolocation process.

Geolocating Images: Where in the world was this taken?

(http://cs231n.stanford.edu/reports/2016/pdfs/404_Report.pdf)

Angela Sy, Cynthia Day

This report page contains a grid of sections and figures. It includes a title, abstract, introduction, related work, methodology, experiments, results, and conclusion. The results section features a large figure showing a confusion matrix and various classification accuracy plots for different flower classes.

Flower Taxonomic Classification using CNNs

(http://cs231n.stanford.edu/reports/2016/pdfs/253_Report.pdf)

Xavier Mignot, Maxwell Siegelman

This report page contains a grid of sections and figures. It includes a title, abstract, introduction, related work, methodology, experiments, results, and conclusion. The results section features a large figure showing a 3D reconstruction of a heart and various plots of left ventricular volume over time.

Convolutional Neural Networks for Estimating Left Ventricular Volume

(http://cs231n.stanford.edu/reports/2016/pdfs/323_Report.pdf)

Ryan Silva, Maksim Korolev

This report page contains a grid of sections and figures. It includes a title, abstract, introduction, related work, methodology, experiments, results, and conclusion. The results section features a large figure showing a 3D reconstruction of a heart and various plots of left ventricular volume over time.

Going Deeper on the Tiny Imagenet Challenge
(http://cs231n.stanford.edu/reports/2016/pdfs/405_Report.pdf)
Andrew Zhai



Tiny ImageNet Visual Recognition Challenge
(http://cs231n.stanford.edu/reports/2016/pdfs/401_Report.pdf)
Anna Shcherbina



Convolutional Neural Network for Traditional Chinese Calligraphy Recognition
(http://cs231n.stanford.edu/reports/2016/pdfs/257_Report.pdf)
Boqi Li



Image Detection Techniques on the Daimler Monocular Pedestrian Detection Benchmark Data
(http://cs231n.stanford.edu/reports/2016/pdfs/288_Report.pdf)
Christopher Ling

Study of Residual Networks for Image Recognition

(http://cs231n.stanford.edu/reports/2016/pdfs/264_Report.pdf)

Mohammad Sadegh Ebrahimi, Hossein Karkeh Abadi

Using Convolutional Neural Network for the Tiny ImageNet Challenge

(http://cs231n.stanford.edu/reports/2016/pdfs/425_Report.pdf)

Jason Ting

Neurofusion: Fusing MEG and EEG Data

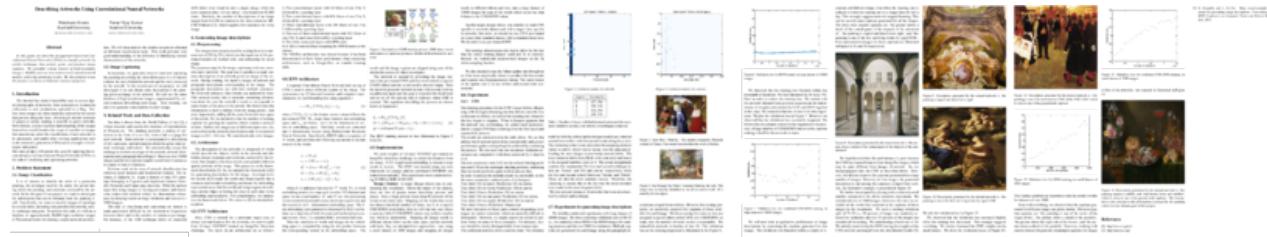
(http://cs231n.stanford.edu/reports/2016/pdfs/321_Report.pdf)

Paul Warren

Virtual art museum

(http://cs231n.stanford.edu/reports/2016/pdfs/200_Report.pdf)

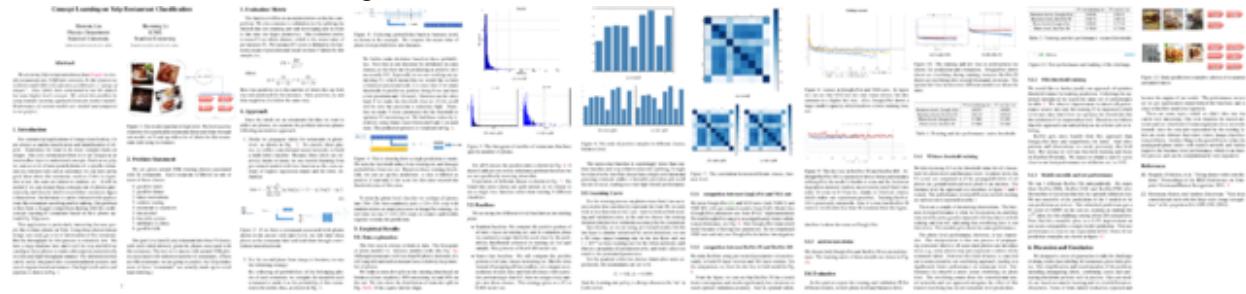
Manikanta Kotaru, Varun Vijay Kumar



Concept Learning on Yelp Restaurant Classification

(http://cs231n.stanford.edu/reports/2016/pdfs/024_Report.pdf)

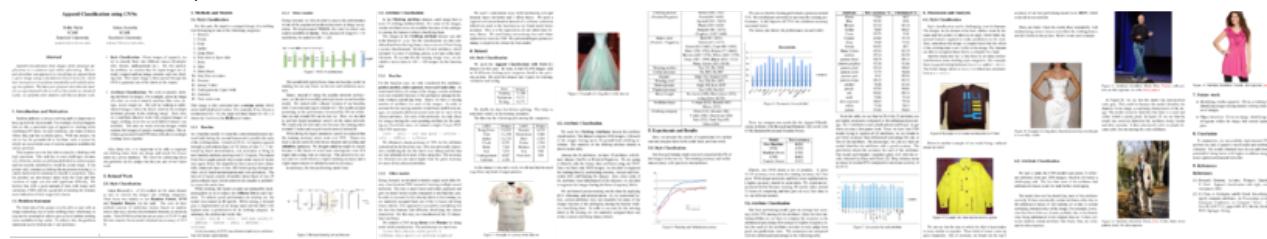
Shenxiu Liu, Haoming Li



Apparel classification using CNNs

(http://cs231n.stanford.edu/reports/2016/pdfs/286_Report.pdf)

Rohit Patki, Suhas Suresha



A network combination of ResNet and GoogLeNet to tackle Tiny ImageNet

(http://cs231n.stanford.edu/reports/2016/pdfs/410_Report.pdf)

William Shen

This section contains a collage of 12 pages from a report. The pages are organized into three columns and four rows. The first column contains pages 1 through 4. The second column contains pages 5 through 8. The third column contains pages 9 through 12. The pages include various diagrams, tables, and images related to the project.

Yelp Restaurant Photo Classification Challenge

(http://cs231n.stanford.edu/reports/2016/pdfs/014_Report.pdf)

Rajarshi Roy

This section contains a collage of 12 pages from a report. The pages are organized into three columns and four rows. The first column contains pages 1 through 4. The second column contains pages 5 through 8. The third column contains pages 9 through 12. The pages include various diagrams, tables, and images related to the project.

Application of Convolved Neural Networks for Pedestrian Detection

(http://cs231n.stanford.edu/reports/2016/pdfs/254_Report.pdf)

Anil Variyar

This section contains a collage of 12 pages from a report. The pages are organized into three columns and four rows. The first column contains pages 1 through 4. The second column contains pages 5 through 8. The third column contains pages 9 through 12. The pages include various diagrams, tables, and images related to the project.

Classifying Restaurants with Yelp Photos

(http://cs231n.stanford.edu/reports/2016/pdfs/021_Report.pdf)

Mike Yu, Emma Marriott, Aaron Zweig

This section contains a collage of 12 pages from a report. The pages are organized into three columns and four rows. The first column contains pages 1 through 4. The second column contains pages 5 through 8. The third column contains pages 9 through 12. The pages include various diagrams, tables, and images related to the project.

Tiny ImageNet Challenge

(http://cs231n.stanford.edu/reports/2016/pdfs/419_Report.pdf)

Vani Khosla

This block contains a collage of several pages from the report. It includes tables of data, plots of training and validation loss, and small images of the dataset. The overall layout is dense and technical, reflecting the nature of the machine learning project.

Identifying the Higgs Boson

(http://cs231n.stanford.edu/reports/2016/pdfs/300_Report.pdf)

Anton Apostolatos, Leonard Bronner

This block contains a collage of several pages from the report. It includes complex plots of particle decay channels and tables of experimental results. The layout is very dense and technical, reflecting the complexity of the physics experiment.

Logo Recognition in Real World Images

(http://cs231n.stanford.edu/reports/2016/pdfs/268_Report.pdf)

Naveen Arivazhagan

This block contains a collage of several pages from the report. It includes tables of data and small images of logos. The layout is dense and technical, reflecting the nature of the computer vision project.